

# Transformer-Based Behaviour Cloning Using PPO Expert Demonstrations

## 1. Abstract

This project shows a full imitation learning pipeline in which a Transformer-based policy learns to balance the cartpole by copying a high-quality PPO expert. The PPO agent learns to act almost perfectly through reinforcement learning. After that, its paths are collected and used to train a Behaviour Cloning (BC) model in a supervised way. The BC model uses a Transformer encoder to find time-based connections between stacked observations. The last model is tested in the environment and can accurately mimic the expert's behaviour, showing that Transformers work well for imitation-based control tasks.

## 2. Introduction to the Algorithm

Imitation Learning (IL) emphasises acquiring decision-making policies through the observation of an expert, rather than through direct interaction with the environment. We use Proximal Policy Optimisation (PPO), a reliable and stable reinforcement learning algorithm, to create expert actions for this project. These expert demonstrations are used to teach a Transformer-based Behaviour Cloning model how to turn a stack of observations into expert actions.

The pipeline is made up of:

1. PPO training to figure out the best way to control a Cart-Pole.
2. A collection of expert demonstrations (state-action pairs).
3. Teaching a Transformer model to copy the expert policy (BC).
4. Testing how well the BC model works and making a video of it in action.

This method is like modern robotics pipelines like ACT, RT-1, and RT-2, where big models learn from expert demonstrations instead of expensive RL interactions.

## 3. Mathematical Interpretation

### 3.1 PPO Expert Training

Cart-Pole is formulated as an MDP:

$$\mathcal{M} = (\mathcal{S}, \mathcal{A}, P, R, \gamma)$$

PPO optimizes the clipped surrogate loss:

$$L^{CLIP}(\theta) = E[\min(r_t(\theta)A_t, \text{clip}(r_t(\theta), 1 - \epsilon, 1 + \epsilon)A_t)]$$

where

$$r_t(\theta) = \frac{\pi_\theta(a_t|s_t)}{\pi_{\theta_{\text{old}}}(a_t|s_t)}$$

This ensures controlled, stable policy updates.

### 3.2 Behaviour Cloning (Supervised Imitation)

Once PPO has converged, we collect dataset:

$$\mathcal{D} = \{(x_t, a_t)\}_{t=1}^N$$

where

$x_t$  is a **stacked observation** and  $a_t$  is the expert action.

BC minimizes cross-entropy:

$$\mathcal{L}(\theta) = - \sum_{t=1}^N \log \pi_\theta(a_t | x_t)$$

### 3.3 Transformer Policy Model

Each stacked observation  $x_t$  is passed through:

#### Embedding

$$h_0 = W_{\text{embed}} x_t$$

#### Self-Attention

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^\top}{\sqrt{d_k}}\right)V$$

#### Action Prediction

$$\pi_\theta(a|x_t) = \text{softmax}(W_{\text{out}} \cdot h_L)$$

The model learns to mimic the PPO expert without interacting with the environment.

## 4. Pseudocode for the Complete Algorithm

### Algorithm: Transformer-Based Behaviour Cloning from PPO Expert

#### Step 1: PPO Expert Training

```

Initialize PPO policy  $\pi_E$ 
for t in 1 ... T:
    s_t ← environment state
    a_t ← sample from  $\pi_E(a | s_t)$ 
    Execute a_t and store transition
    Update  $\pi_E$  using PPO clipped objective
Save  $\pi_E$ 
```

## Step 2: Dataset Collection

```
Initialize dataset D = {}  
for each episode:  
    Reset env → s  
    Initialize a stack of last K observations  
    while episode not done:  
        a = π_E(s)  
        x = stacked observation  
        Add (x, a) to D  
        s ← next state  
    Save D
```

## Step 3: Train Transformer for Behavior Cloning

```
Load dataset D & Normalize observations  
Initialize Transformer policy π_θ  
for epoch = 1 to E:  
    for each mini-batch (x_b, a_b):  
        logits = π_θ(x_b)  
        loss = CrossEntropy(logits, a_b)  
        Backpropagate loss and update θ  
    Save π_θ
```

## Step 4: Evaluate and Record

```
Load model π_θ  
for each evaluation episode:  
    Reset env  
    while not done:  
        x ← stacked observation  
        a_pred = argmax π_θ(a | x)  
        Step environment and record frame  
    Save frames as MP4/GIF
```

## 5. Output in Terms of Iterations

This section interprets the detailed training logs obtained during execution, particularly the **PPO expert's iteration-by-iteration progress**.

The file shows around **75 iterations**, each corresponding to **2048 environment steps**, producing approximately **150,000 PPO timesteps** overall.

### Key Observations from Iteration Logs

**Iterations increase steadily from 1 to 75**, showing consistent PPO training progression. At each iteration, PPO prints metrics such as:

- `approx_kl` → how much the new policy diverges from old policy
- `clip_fraction` → how often the objective is clipped
- `entropy_loss` → how random the policy is
- `value_loss` → critic network error
- `explained_variance` → how well value function predicts returns

Over time:

- **approx\_kl decreases** → policy becomes stable
- **entropy\_loss decreases** → policy becomes deterministic (typical for CartPole)
- **explained\_variance increases** → critic becomes accurate
- **value\_loss drops** → critic predictions improve

These results show **convergence** of PPO before collecting demonstrations.

### **BC Training Output (Iterations as Epochs)**

BC training shows:

Epoch	Loss
1	0.08368
10	0.02442
20	0.01648
30	0.01506
40	0.01323
50	0.01148
60	0.01163

This shows:

- Sharp improvement in the first 10 epochs
- Loss plateauing after epoch 40
- Model has **successfully learned** the expert policy

### **Evaluation Output (Iterations as Episodes)**

During evaluation:

```
BC eval rewards: [500.0, 500.0]
```